

Abstract

1 Introduction

Mobile phone surveys are an increasingly popular method for tracking household food security that is employed by aid organizations like the World Food Programme. Improving sampling capabilities and enabling real-time tracking of consumption and coping metrics, these surveys are conducted daily in an array of countries. Drawn with a random sampling strategy to select participants, samples do not necessarily overlap day to day, but are conducted at a frequency that often results in temporal correlation in the collected responses that hinders inference when using linear regression. While there already exist methods for analyzing data in which samples are surveyed repeatedly or aggregated at the time-step level, models that account for individual-level effects and population-level shocks are scarce. Those that do exist model temporal fluctuations with a latent time series that is recursively estimated and used to update parameter estimates via an EM algorithm, but these methods do not properly account for how cross-sectional noise impacts analysis. We correct this shortcoming by incorporating methods for analyzing data that are the result of multiple noisy observations of each element of a time series, known as an ARMA-Noise model.

We go on to establish that the use of an EM algorithm that factors the likelihood equation by conditioning on an estimate of a latent ARMA process incorrectly assumes that the resulting deviations will have a covariance structure conducive to standard fixed effect regression. We therefore replace the EM algorithm with one that recursively estimates variance parameters with an ARMA-Noise model and uses them to calculate GLS estimators for regression parameters. Confidence regions for time series parameters are generated via the Delta Method due to their latent nature. The respective implications of varying orders of the latent time ARMA process inform a decision rule for model selection.

While the newly incorporated method for fitting time series parameters corrects the bias that previously existed, it imposes a constant sample size at the time-step level, an unrealistic requirement in real-life surveying contexts. We therefore explored the extent to which ARMA-Noise modeling methods could be generalized to accommodate varying sample sizes.

With the proposed method fully developed, we test on simulated data and illustrate how model selection might be performed on a realization from a restricted model. We also include a case study on household survey data from Honduras that shows the method's potential to decorrelate residuals. The results of our analysis support the method as a viable option for modeling RCS survey data related to food security and other contexts.

2 Background

Repeated cross-sectional (RCS) data has long been the subject of analysis. When cross-sections are taken with high frequency, the possibility of temporal correlation arises. While panel data consist of a single sample being observed repeatedly, RCS data do not necessarily have any repeated observations, ruling out the treatment of the data as a multivariate time series. It is often instead the population mean that is treated as a time series, with individual observations modeled as a realization of that mean with some added noise.

Among the earliest methods designed for RCS data was that proposed by Scott and Smith [6], in which dependent variables are aggregated at the time step level and modeled as a time series with added noise:

$$y_t = \theta_t + e_t, \quad \text{Var}(e_t) = s_t^2 \quad (2.1)$$

In this model, surveys are conducted at regular time intervals and θ_t represents a population mean changing over time. The analytical method proposed therein introduces different possibilities for the covariance structures of $\{\theta_t\}$ and $\{e_t\}$, but is ultimately concerned with the use of all previously observed aggregates $\{y_t\}_{t=1}^T$ to estimate θ_T . As our objective is to improve the estimation and of covariate effects and their standard errors, we must look to more highly parameterized models.

Such a model was developed by DiPrete and Grusky [3], in which added complexity and flexibility is added to a standard fixed effects model by having the effects of individual-level variables depend on macro-level variables:

$$y_{ti} = X_{ti}\beta_t + \varepsilon_{ti}, \quad \beta_t = \mathbf{Z}_t\gamma, \quad \varepsilon_{ti} \sim \text{WN}(0, \sigma^2) \quad (2.2)$$

Such a model requires the presence of macro-level variables that vary over time and explain changing effects of individual-level variables. While this formulation has a broad range of applications, it is only appropriate in cases where temporal correlation can be assumed to be a function of known time-varying predictors, rather than an independent random process.

There also exist aggregate methods that involve calculating averages of independent and dependent variables at each time step and analyze these averages as a time series with covariates. Such a model is an example of a state-space model that can be modeled jointly via Kalman filtering [2]. In the regression setting, aggregates would be needed for both response and explanatory variables, fitting the following model:

$$\bar{\mathbf{y}} = (n^{-1} \mathbf{X}'_t \mathbf{X}_f) \beta + \mathbf{u}, \quad \mathbf{u} \sim \text{ARMA}(p, q)$$

While this method is effective at estimating population-level fixed effects like intercepts and slopes, it is limited when it comes to the estimation of individual-level effects, especially if indicator variables vary little in their distribution at each time step. This is a highly plausible scenario when it comes to sampling methods for surveys, as researchers interested in the effect of a certain covariate

may very well aim to be consistent in the proportions of time step samples that have this feature.

One model that has been developed to analyze RCS data with autocorrelated errors is the Autoregressive Random Effects (ARE) Model [5]. Treating observations as the sum of a fixed effects model and an AR(1) time series, ARE models have two levels: a fixed-effects level and a time series level.

The method employed to fit such a model is a recursive process that splits the log likelihood function of the model into two terms by conditioning on the latent time series $\mathbf{u}$. An EM algorithm is then employed for parameter estimation.

This conditional splitting is convenient, as $\mathbf{u} \sim N(\mathbf{0}, \mathbf{\Gamma})$ and $\mathbf{y}|\mathbf{u} \sim N(X_f\beta + \mathbf{u}, \sigma_\varepsilon^2 I_N)$. Furthermore, the choice of imposing an AR(1) structure on $\mathbf{u}$ lends itself to a concise representation of the determinant and inverse of $\mathbf{\Gamma}$, the covariance matrix of $\mathbf{u}$.

In the absence of an explicit observation of $\mathbf{u}$, its conditional expectation is calculated with current parameter estimates, then used to update those estimates. This process is repeated until convergence criteria are met.

There are two flaws in this methodology. The first is that fixed effect parameter estimates are updated with the assumption that $\mathbf{y} - \hat{\mathbf{u}}$ is a realization of a fixed effects model, when its covariance structure is altered by the error in estimating $\mathbf{u}$. The second problem stems from the effect of added individual noise on an AR(1) process, the sum of which forms an ARMA-Noise Process [4].

Let's assume the best case scenario for level one of the model, namely the perfect estimation of fixed effect parameters ($\hat{\beta} = \beta$). In this unlikely event, we are able to extract data consisting of multiple noisy observations of each element of an AR(1) process

With no prior information about the latent time series, the best estimates for the elements of $\mathbf{u}$ are the averages at each time step. The added noise turns an originally AR(1) process and turns it into an ARMA(1,1) process. Failing to add a moving average parameter necessarily introduces bias to the estimation of the time series parameters.

From Box (1976) [1], we have the following theorem:

If $\{y_t\}$ follows an ARMA(p_1, q_1) model, $\{a_t\}$ follows an ARMA(p_2, q_2) and the two series are independent, then their sum $w_t = y_t + a_t$ follows an ARMA(p, q) model in which $p \leq p_1 + p_2$ and $q \leq \max(p_1 + q_2, p_2 + q_1)$.

Because $\{a_t\}$ is white noise, $p_2 = q_2 = 0$. As $p_1 = 1$ and $q_1 = 0$, we can deduce that $\{w_t\}$ is an ARMA(p, q) process where $p \leq 1$ and $q \leq 1$. Because $\gamma_{\hat{\mathbf{u}}}(h+1)/\gamma_{\hat{\mathbf{u}}}(h) = \phi$ for $h > 0$, we know that $p = 1$. The ACVF of $\hat{\mathbf{u}}$ is therefore incompatible with that of an AR(1) process, so it must be ARMA(1,1).

As the best-case scenario of estimating and extracting the latent time series will still yield temporal correlation in the data, albeit reduced, we must turn to a more nuanced treatment of residual covariance. A natural choice is the GLS estimator of β , which we will see challenges the necessity of using $\hat{\mathbf{u}}$ in the estimation of β at all.

When framed in the context of repeated surveys, the requirement that the sample size at each time step is constant is unrealistic. Given that real life

survey data is rarely so uniform, generalizing the method to allow for varying daily sample size is a natural goal. In examining the method as it stands, the condition of fixed sample size is only necessary in the fitting of ARMA-Noise models, as $\mathbf{w} = \mathbf{u} + \mathbf{a}$ is only an ARMA process if $\{a_t = n_t^{-1} \sum_{i=1}^{n_t}\}$ is white noise, something that at first glance is violated by varying n_t . However, we will now show how this method can be adapted to allow for varying sample sizes, so long as certain conditions are met.

3 Method

Having identified what parts of existing methods we want to utilize as well as those that we wish to change or generalize, we are ready to define our model and accompanying algorithm, *REgression Models for Autocorrelated Repeated Cross-Sections*, or REMARCS:

3.1 Model and Algorithm

We define the vector of observations $\mathbf{y}$ as the realization of a model defined as the sum of fixed effects, a latent ARMA(1,1) process, and i.i.d. noise (*n.b. sample sizes are currently constant across time steps*).

$$y_{ti} = \mathbf{x}_{ti}\beta + u_t + \epsilon_{ti}, \quad i = 1, \dots, n \quad \epsilon_{ti} \sim \text{i.i.d. } N(0, \sigma_\epsilon^2) \quad (3.1)$$

$$u_t - \phi u_{t-1} = \eta_t + \theta \eta_{t-1}, \quad t = 1, \dots, T \quad \eta_t \sim N(0, \sigma_\eta^2) \quad (3.2)$$

The algorithm we propose is an iterative process applying the estimates of parameters from one level to update those of the other. The process is as follows:

- 1) Initialize with $\hat{\beta} = (X_f' X_f)^{-1} X_f' \mathbf{y}$
- 2) Estimate $(X_t \mathbf{u} + \varepsilon)$ with $(\mathbf{y} - X_f \hat{\beta})$ and fit ARMA-Noise model to obtain estimates $(\hat{\phi}, \hat{\theta}, \hat{\sigma}_\eta^2, \hat{\sigma}_\epsilon^2)$
- 3) Calculate $\hat{\Omega}$ and $l(\hat{\xi})$
- 4) Refit fixed effect parameter estimates by calculating GLS estimate $\hat{\beta}_{GLS} = (X_f' \hat{\Omega}^{-1} X_f)^{-1} X_f' \hat{\Omega}^{-1} \mathbf{y}$
- 5) Repeat Steps 2-4 until convergence criteria are met
- 6) Generate confidence regions. If any parameters are insignificant, restrict model and repeat algorithm

MAKE SURE ALL NOTATION HAS BEEN DEFINED!!!!




Figure 1: Flowchart of algorithm steps for estimating model parameters

3.2 Evaluating Convergence

Instead of splitting the likelihood by conditioning on $\hat{\mathbf{u}}$, doing so with $\bar{\mathbf{y}}$ instead of has the advantage that likelihoods are based on a constant statistic.

Let's recall the log likelihood $l(\xi)$:

$$l(\xi) = -\frac{N}{2}\ln(2\pi) - \frac{1}{2}\ln|\Omega| - \frac{1}{2}(\mathbf{y} - X_f\beta)' \Omega^{-1}(\mathbf{y} - X_f\beta)$$

We now condition on $\bar{\mathbf{y}} := n^{-1}X'_t\mathbf{y}$, which has the following distribution.

$$\bar{\mathbf{y}} \sim N(\mu_{\bar{\mathbf{y}}}, \Sigma_{\bar{\mathbf{y}}}), \quad \mu_{\bar{\mathbf{y}}} = n^{-1}X'_tX_f\beta, \quad \Sigma_{\bar{\mathbf{y}}} = \Gamma + n^{-1}\sigma_\varepsilon^2 I_T$$

Defining ξ to be the vector of parameters $(\beta, \phi, \theta, \sigma_\eta^2, \sigma_\varepsilon^2)$, we can then define the log likelihood function:

$$l_t(\xi) = \ln[f(\bar{\mathbf{y}}|\xi)] = -\frac{T}{2}\ln(2\pi) - \frac{1}{2}\ln|\Sigma_{\bar{\mathbf{y}}}| - \frac{1}{2}(\bar{\mathbf{y}} - \mu_{\bar{\mathbf{y}}})' \Sigma_{\bar{\mathbf{y}}}^{-1}(\bar{\mathbf{y}} - \mu_{\bar{\mathbf{y}}})$$

If we then condition $\mathbf{y}$ on $\bar{\mathbf{y}}$, the only randomness left is in the vector of deviations, which are independent of $\bar{\mathbf{y}}$:

$$\mathbf{y} - X_t\bar{\mathbf{y}} = \boldsymbol{\varepsilon} - X_t\mathbf{a}$$

This is a vector with mean $\mathbf{0}$ and variance:

$$\text{Var}(\boldsymbol{\varepsilon} - X_t\mathbf{a}) = \sigma_\varepsilon^2(I_N - n^{-1}X_tX'_t) \quad (3.3)$$

We now derive the second part of our split likelihood. Because of the rank-deficient nature of the matrix Equation 3.3, we can simplify the log likelihood

$$l_i(\xi) = \ln f(\mathbf{y} | \bar{\mathbf{y}}, \sigma_\varepsilon^2) = -\frac{N-T}{2}\ln(2\pi\sigma_\varepsilon^2) - \frac{1}{2\sigma_\varepsilon^2} \sum_{t=1}^T \sum_{i=1}^n (y_{it} - \bar{y}_t)^2, \quad (3.4)$$

With this we can express the full likelihood as $l(\xi) = l_t(\xi) + l_i(\sigma_\varepsilon^2)$. However, since $\bar{\mathbf{y}}_t$ is independent of our model selection and parameter estimates, l_i is going to be relatively stable. We therefore pair $\xi \in \mathbb{R}^k$ and $l_t(\xi)$ to calculate the Akaike Information Criterion (AIC):

$$AIC = -2k - 2l_t(\xi)$$

Calculating this weighted likelihood is how we will compare the performance of models that include different subsets of parameters.

3.3 Inference

Once parameter estimates have converged, confidence intervals are constructed.

Fixed effect estimates have covariance matrix:

$$\text{Var}(\hat{\beta}_{GLS}) = (X_f' \Sigma_{\mathbf{y}}^{-1} X_f)^{-1} \quad (3.5)$$

This covariance matrix is estimated with final parameter estimates and then used to construct confidence regions. Just as this method of building confidence regions for β in some sense assumes that parameters $\phi, \theta, \sigma_\eta^2, \sigma_\varepsilon^2$ have been perfectly estimated, confidence intervals for those same parameters makes an analogous assumption.

If β is estimated perfectly, then $\mathbf{y} - X_f \hat{\beta} = X_t \mathbf{u} + \varepsilon =: \mathbf{z}$, and time-step averages $w_t = \sum_{i=1}^n z_{ti}$ are equal to the sum $u_t + a_t$, with $a_t \stackrel{i.i.d.}{\sim} N(0, \sigma_\varepsilon^2/n)$. According to Box and Jenkins [1], if $\mathbf{u}$ is an ARMA(p_u, q_u) process with $p_u, q_u \leq 1$, then $\mathbf{w}$ is an ARMA(p_w, q_w) process, with $p_w, q_w \leq 1$.

What's most important to note is that if $\mathbf{u}$ is parameterized by $\phi, \theta, \sigma_\eta^2$, then $\mathbf{w}$ is parameterized by $\phi, \theta_b, \sigma_b^2$. According to [2], the Maximum Likelihood estimates for $\phi, \theta_b, \sigma_b^2$, which can be obtained using an R package like *arima*, have the following asymptotic distribution (when normalized):

$$\sqrt{T} \left[\begin{pmatrix} \hat{\phi} \\ \hat{\theta}_b \\ \hat{\sigma}_b^2 \end{pmatrix} - \begin{pmatrix} \phi \\ \theta_b \\ \sigma_b^2 \end{pmatrix} \right] \xrightarrow{d} N(\mathbf{0}, \Sigma_b)$$

$$\Sigma_b = \begin{bmatrix} \frac{(1-\phi^2)(1+\phi\theta_b)^2}{(\phi+\theta_b)^2} & \frac{-(1-\theta_b^2)(1-\phi^2)(1+\phi\theta_b)}{(\phi+\theta_b)^2} & 0 \\ \frac{-(1-\theta_b^2)(1-\phi^2)(1+\phi\theta_b)}{(\phi+\theta_b)^2} & \frac{(1-\theta_b^2)(1+\phi\theta_b)^2}{(\phi+\theta_b)^2} & 0 \\ 0 & 0 & 2\sigma_b^4 \end{bmatrix}$$

The parameter estimate $\hat{\sigma}_\varepsilon^2$ and its normalized asymptotic variance are given by:

$$\hat{\sigma}_\varepsilon^2 = (N - T)^{-1} \sum_{t=1}^T \sum_{i=1}^{n_t} (w_{ti} - \bar{w}_t.)^2 \quad (3.6)$$

$$\sqrt{T}(\hat{\sigma}_\varepsilon^2 - \sigma_\varepsilon^2) \xrightarrow{d} N(0, \frac{2\sigma_\varepsilon^4}{n-1})$$

Because $y_t. \perp S_t^2$ for all t , we know that $\hat{\sigma}_\varepsilon^2 \perp (\hat{\phi}, \hat{\theta}_b, \hat{\sigma}_\eta^2)$ and can give the following asymptotic distribution for “observable” parameters $\tau_o = (\phi, \theta_b, \sigma_b^2, \sigma_\varepsilon^2)$:

$$\sqrt{T}(\hat{\tau}_o - \tau_o) \xrightarrow{d} N(\mathbf{0}, \Sigma_o), \quad \Sigma_o = \begin{bmatrix} \Sigma_b & \mathbf{0} \\ \mathbf{0} & \frac{2\sigma_\varepsilon^2}{n-1} \end{bmatrix} \quad (3.7)$$

Equipped with the above estimates, Wong and Miller detail how to estimate “latent” parameters $\tau_l = (\theta, \sigma_\eta^2)$ by relating the autocovariance functions of $\mathbf{u}$ and $\mathbf{w}$ through the system of equations $H(\theta, \sigma_\eta^2)$:

PASTE H EQUATION HERE!!!!!!

$$\tau_0 = \begin{pmatrix} \phi \\ \theta_b \\ \sigma_b^2 \\ \sigma_\varepsilon^2 \end{pmatrix} \quad \tau_1 = \begin{pmatrix} \phi_* \\ \sigma_{\varepsilon*}^2 \\ \theta \\ \sigma_\eta^2 \end{pmatrix} \quad (3.8)$$

Estimates for τ_l are then given by the roots of H . It is important to note that estimates for τ_o can produce a system of equations that has no real roots, especially for relatively low values of T . In this event we use numerical methods to solve for parameter estimates $\hat{\theta}, \hat{\sigma}_\eta^2$ that minimize $\|H\|^2$.

Were we to have an explicit function $(\hat{\phi}, \hat{\theta}, \hat{\sigma}_\eta^2) = F(\hat{\tau}_o)$ providing estimates for τ_l , we would directly apply the Delta Method to obtain the asymptotic distribution of our vector of parameter estimates:

$$\sqrt{T}(F(\hat{\tau}_o) - \tau_l) \xrightarrow{d} N(\mathbf{0}, G\Sigma_o G'), \quad G = \frac{\partial F}{\partial(\tau_o)} \quad (3.9)$$

Because we're using H , a function whose roots are our estimates, we will adapt 3.9 using the Implicit Function Theorem:

$$G = - \left(\frac{\partial H_f}{\partial \tau_1} \right)^{-1} \left(\frac{\partial H}{\partial \tau_0} \right) \quad (3.10)$$

Because we're ultimately interested in the normalized asymptotic distribution of parameter estimates $(\hat{\phi}, \hat{\theta}, \hat{\sigma}_\eta^2, \hat{\sigma}_\varepsilon^2)$, we expand H to include some identity functions.

Our procedure is to fit $\hat{\tau}_o$, solve for $\hat{\tau}_l$, and then plug in all estimates into Equations 3.9 and 3.10. If any parameter estimates, observed or latent, are deemed insignificant, the relevant parameter can be restricted to 0 and the procedure adapted accordingly. This decision procedure is detailed next.

3.4 Model Selection

Before outlining the methodology for model selection, we motivate its design by exploring the implications of certain parameter combinations in an ARMA-Noise model $\mathbf{w} = \mathbf{u} + \mathbf{a}$.

Parameter Restriction	Order of $\mathbf{u}$	Order of $\mathbf{w}$
$\theta = 0$	AR(1)	ARMA(1,1)
$\phi = 0$	MA(1)	MA(1)
$\frac{\theta}{\phi} = \frac{\sigma_\varepsilon^2}{n\sigma_\eta^2}$	ARMA(1,1)	AR(1)
$\phi = \theta = 0$	WN	WN
Else	ARMA(1,1)	ARMA(1,1)

Table 1: Implications of Various Restrictions of Time Series Parameters

From these relations we derive a decision rule for model selection:




Figure 2: Decision tree for selecting parameters to include in ARMA-Noise model

3.5 Generalizing to Varying Sample Size

Having detailed the steps involved in fitting and performing inference on model parameter estimates, we now turn to the condition required for fitting an ARMA-Noise model, namely that sample size is constant at each time step ($n_t = n \forall t$). This condition is included so that the time series $\{a_t\}$ is Gaussian white noise, and $\mathbf{w}$ is therefore still an ARMA process.

Given that RCS data is rarely of a fixed sample size, we now explore the extent to which this condition can be relaxed whilst $\mathbf{a}$ remains white noise. We will now treat $\{n_t\}$ as the realizations of random variables $\{N_t\}_{t=1}^T$. What we aim to establish are the conditions under which the random variable $A_t = N_t^{-1} \sum_{i=1}^{N_t} \varepsilon_{ti}$ is still white noise.

Claim:

Let $\{\varepsilon_{ti}\}$ be i.i.d. random variables with

$$\mathbb{E}(\varepsilon_{ti}) = 0, \quad \text{Var}(\varepsilon_{ti}) = \sigma_\varepsilon^2 < \infty,$$

Let $\{N_t\}$ be a stochastic process such that

$$P(N_t \in \mathbb{N} \setminus \{0\}) = 1 \quad \mathbb{E}(N_t^{-1}) \quad \text{for all } t,$$

Define

$$A_t := \frac{1}{N_t} \sum_{i=1}^{N_t} \varepsilon_{ti}.$$

Then $\{A_t\}$ is white noise with $\sigma_a^2 = \sigma_\varepsilon^2 \mathbb{E}(N_t^{-1})$

Proof:

First, let $\{N_t\}$ be i.i.d. random variables with a sample space that is a subset of the integers strictly greater than 0. The conditional random variable $A_t|N_t$ then has variance σ_ε^2/N_t .

To find the marginal variance $\text{Var}(A_t)$, we use the Law of Total Variation:

$$\text{Var}(A_t) = \mathbb{E}(\text{Var}(A_t|N_t)) + \text{Var}(\mathbb{E}(A_t|N_t)) \quad (3.11)$$

Now, $\text{Var}(A_t|N_t) = \sigma_\varepsilon^2 N_t^{-1}$, and therefore has an expectation of $\sigma_\varepsilon^2 \mathbb{E}(N_t^{-1})$, which is well defined given that the sample space of N_t is a subset of the integers strictly greater than zero. At the same time, the expectation of A is 0 regardless of the value of N , meaning the second term in Equation 3.11 is zero, leaving only $\text{Var}(A) = \sigma_\varepsilon^2 \mathbb{E}(N^{-1})$.

Proof:

$$N \geq 1 \quad \text{a.s.}$$

$$\Rightarrow 0 < N^{-1} \leq 1 \quad \text{a.s.}$$

$$\Rightarrow 0 < \mathbb{E}(N^{-1}) \leq 1 \quad \text{a.s.}$$

We can quickly verify that A_t is zero mean with the Law of Total Expectation:

$$\mathbb{E}(A_t) = \mathbb{E}(\mathbb{E}(A_t|N_t)) = \mathbb{E}(0) = 0$$

Now we've already established $\text{Var}(A_t)$ to be independent of t , so we now verify that $\text{Cov}(A_r, A_s) = 0$ for $r \neq s$. The Law of Total Covariance states:

$$\text{Cov}(A_r, A_s) = \mathbb{E}(\text{Cov}(A_r, A_s|N_r, N_s)) + \text{Cov}(\mathbb{E}(A_r|N_r), \mathbb{E}(A_s|N_s))$$

Given that A_r and A_s are functions of independent random vectors $\varepsilon_r \in \mathbb{R}^{n_r}$ and $\varepsilon_s \in \mathbb{R}^{n_s}$, $\text{Cov}(A_r, A_s|N_r, N_s)$ must be zero. Having already established that both variables are zero-mean, the above equation reduces to:

$$\text{Cov}(A_r, A_s) = \mathbb{E}(0) + \text{Cov}(0, 0) = 0$$

□

Having shown that $\{A_t\}$ is still white noise, it now remains to estimate its variance. Though the sample sizes now vary, our estimate for σ_ε^2 hasn't changed (see Equation 3.6).

Without requiring the distribution of N_t , σ_a^2 can be estimated with:

$$\hat{\sigma}_a^2 = \frac{\hat{\sigma}_\varepsilon^2}{T} \sum_{t=1}^T \frac{1}{N_t} \quad (3.12)$$

With $\hat{\sigma}_\varepsilon^2$, $\hat{\sigma}_a^2$ and $\hat{\text{Var}}(\hat{\sigma}_\varepsilon^2)$ defined as above, the algorithm proposed by Wong and Miller can be once again applied, with necessary adjustments made to Σ_o and H .

The takeaway from all this is that a constant expected inverse of $\{N_t\}$ is the critical piece for $\{A_t\}$ to still be white noise.

So long as the set of inverse sample sizes $\{n_t^{-1}\}$ is conceivably of constant expectation, the ARMA-Noise method is applicable.

COMMENT ON ASSUMPTION OF INDEPENDENCE IN CALCULATING VARIANCE OF ESTIMATE AND IMPLICATIONS FOR STRICTLY STATIONARY SAMPLE SIZES AND NON-CONSTANT MEAN

3.6 Extending to Mixed Effects Models

INITIAL FIT OF MIXED EFFECTS MODEL

UPDATE TO GLS ESTIMATOR FOR BETA

CONDITIONAL EXPECTATION OF RANDOM INTERCEPTS

ESTIMATE FOR RANDOM EFFECT VARIANCES

4 Results

- SIMULATE FIXED/MIXED EFFECT REMARCS DATA
- FIT WITH FE/ME MODEL
- FIT WITH TS WITH COVARIATES
- FIT WITH MODUGNO
- MODUGNO INSTABILITY
- FIT WITH REMARCS
- HEATMAPS
- MODEL SELECTION

A realization of an ARE model is displayed in Figure 3 with the following parameters:

Parameter	β_0	β_t	β_1	σ_ε^2	ϕ	σ_η^2
Value	40	-0.1	2	1	0.9	1

Table 2: Parameter values for simulated ARE model




Figure 3: Realization of ARE model with parameters defined in Table 2

4.1 Simulations

In order to evaluate the efficacy of the methods proposed above, individual components of the REMARCS methodology and the algorithm in its entirety were tested on simulated RCS data.

We consider fitting this model by aggregating the data at each time point and treating the averages as a time series with covariates. 100 simulations were conducted with the following parameters:

$$\beta = (43, -.001, .25), \quad \sigma_\varepsilon^2 = 100, \quad \phi = \theta = .5, \quad \sigma_\eta^2 = 1$$

Some sample rows of simulated data are provided below:

t	x	y
1.00	0	48.14
1.00	0	25.17
1.00	0	42.60
1.00	1	68.78
1.00	0	37.36
1.00	0	30.96

After aggregating at each time step, we obtain the following condensed dataset, from which a time series with covariates is then fit:

t	n	x	$\bar{y}$
1.00	94	0.37	44.88
2.00	84	0.50	43.24
3.00	96	0.47	43.49
4.00	107	0.52	42.68
5.00	89	0.47	42.45
6.00	99	0.49	41.87

Note that the x column of each time step sample has a value close to 0.5, with some variability. If we then fit a time series with covariates to the condensed dataset, we obtain results with empirical means displayed in the table below. The “Significance Frequency” column displays the fraction of the hundred simulations in which the corresponding parameter estimate was statistically significant at the $\alpha = .05$ level.

Parameter	True Value	Estimate Mean	Sig. Freq.
phi	0.500	0.494	1.000
theta	0.500	0.001	0.040
Intercept	43.000	43.017	1.000
Slope	-0.001	-0.001	0.860
Covariate	0.250	0.216	0.110

Note that the effect of the dummy indicator variable is deemed significant only 11% of the time. This example highlights the shortcomings of fitting a time series with covariates. Also of note is how skewed the sample mean of $\hat{\theta}$ is, in particular how close to zero it is. We will see later why this is actually expected. For now, we use these results to motivate the need for a different approach. With

that need in mind, we next review some of the key developments in modeling repeated cross-sectional data.




Figure 4: Rates at which point estimates for τ_o fall within confidence regions based on true values

We begin with simulations of an ARMA(1,1)-Noise Model $\mathbf{y} = X_t \mathbf{u} + \varepsilon$ parameterized by $\phi, \theta, \sigma_\eta^2, \sigma_\varepsilon^2$. Empirical confidence levels for each value-pair (ϕ, θ) are based on a hundred simulations with $\sigma_\eta^2 = \sigma_\varepsilon^2 = 1$, and $n_t = 2$. We first verify that the time series $\mathbf{w}$ defined by $w_t = \bar{\mathbf{y}}_t$, with MA parameter θ_b and innovation variance σ_b^2 , produces time series estimates with normalized covariance close to Σ_{ϕ, θ_b} . For each set of parameters, we calculate the true values of θ_b and σ_b^2 . We then fit $(\hat{\phi}, \hat{\theta}_b)$ with $\bar{\mathbf{w}}$ and construct confidence regions centered around those estimates, counting the rates at which these confidence regions capture the true parameter values. These empirical confidences are displayed in Figure 4.





Figure 5: Rates at which point estimates for τ_l fall within confidence regions based on sample estimates

Confidence regions calculated with $\alpha = .05$ capture true values at rates close to the predicted confidence of 95%. As we are ultimately interested in estimating the parameters $\phi, \theta, \sigma_\eta^2$, we now see how well confidence regions based on applying the Delta Method to Σ_o perform. The results in Figure 5 were generated by simulating realizations of an ARMA(1,1)-Noise, estimating time series parameters ϕ and θ , then using the Delta Method to construct 95% confidence intervals.

For both observable and latent parameters we can see performance dipped for parameter combinations (ϕ, θ) close to the $\phi + \theta = 0$ line for time series of length $T = 1000$. However, we can see that performance improved for $T = 10,000$, suggesting that parameter estimates for ARMA(1,1) processes that resemble white noise take longer to converge to their asymptotic distributions.

In order to assess the efficacy of this method, simulations were conducted with varying combinations of ARMA parameters (ϕ, θ) . In Figure 6, we can see the proportion of 100 simulations that yielded confidence regions successfully capturing $\beta = (43, -0.001, 0.25)$ (left) and $(\phi, \theta, \sigma_\eta^2)$ (right). One observation we can make is that the estimation of time series parameters appears to be hindered close to the line defined by $\phi + \theta = 0$. This under-performance is mitigated as T grows, indicating that time series more closely resembling white noise will require larger values of T for accurate estimation.



(a) Fixed Effect Parameter Estimation

(b) Time Series Parameter Estimation

Figure 6: Empirical confidences with which model parameters fall within confidence regions generated after model convergence

4.2 Development of model parameter estimates

Now that we've seen some aggregate results displaying the efficacy of this method under different parameters, we highlight some individual realizations of specific parameter combinations and observe how the parameter estimates evolve as the algorithm progresses. We start with a single simulation of the most general case of our model, one in which all parameters are nonzero. Exact values are given below:

T	sims	ϕ	θ	σ_η^2	σ_ε^2	β_0	β_1	β_2
1000	1	0.9	-0.6	1	10	40	-0.001	2

Table 3: Parameter values for simulated REMARCS model with all nonzero parameters

We begin by displaying the individual data, along with the eventual fitted values.




Figure 7: A realization of a REMARCS model before and after fitting

Once the convergence criterion of $|l(\hat{\xi}^{(k+1)}) - l(\hat{\xi}^{(k)})| < .001$ is satisfied, we take a look at the significance of “observable” time series parameter estimates $(\hat{\phi}, \hat{\theta}_b)$. Seeing that both $\hat{\phi}$ and $\hat{\theta}_b$ are significant in the left panel of Figure 8, we proceed to fit estimates for θ and σ_η^2 using H in Equation ???. The right panel of the same figure shows the resulting confidence region after applying the Delta Method, and we can see that the true parameters are successfully captured.




Figure 8: Confidence regions for time series parameter estimates of latent ARMA(1,1) process

In the table below we can see how properly accounting for temporal corre-

lation helps update parameter estimates:

Parameter	β_0	β_1	β_2	σ_ε^2	ϕ	θ	σ_η^2
True Value	40.0000	-0.0010	2.0000	10.0000	0.9000	-0.6000	1.0000
Initial estimate	40.4306	-0.0016	1.9526	11.5826	0.0000	0.0000	0.0000
Final estimate	40.4352	-0.0017	1.9857	9.8552	0.9296	-0.5955	0.9482

Table 4: Parameter estimates for REMARCS model with all nonzero parameters, before and after convergence

We now look at data simulated from a restricted model with the following parameters:

T	sims	ϕ	θ	σ_η^2	σ_ε^2	β_0	β_1	β_2
1000	1	0.8	0	10	10	-10	0.001	1

Table 5: Parameter values for simulated REMARCS model with a latent AR(1) process




Figure 9: Observations and fitted values from a REMARCS model with a latent AR(1) time series

Parameter	β_0	β_1	β_2	σ_ε^2	ϕ	θ	σ_η^2
True Value	-10.0000	0.0010	1.0000	10.0000	0.8000	0.0000	10.0000
Initial estimate	-10.2365	0.0011	1.1015	38.6200	0.0000	0.0000	0.0000
Final estimate	-10.0172	0.0007	0.9519	9.8354	0.8130	-0.0673	10.7934

Table 6: Parameter estimates REMARCS model with latent AR(1) Time Series, before and after convergence

After having applied the general method to observations and the convergence of parameter estimates, we now perform inference on $(\hat{\phi}, \hat{\theta}_b)$. Seeing that both estimates are significant in the left side of Figure 9, we proceed as before with the most general form of H .





Figure 10: Confidence regions for time series parameter estimates of latent ARMA(1,1) process

We then see in the right panel of Figure 9 that the resulting estimate for θ is insignificant, prompting us to repeat the estimation of σ_η^2 by using H_θ in Equation ???. While the standard error for $\hat{\phi}$ is unaffected, the restriction of $\hat{\theta}$ will affect $\hat{\sigma}_\eta^2$. To complete this analysis, the algorithm would be restarted, now with the restriction $\hat{\theta} = 0$.

Figure 11: Confidence interval for ϕ after restricting model

We now restrict $\phi = 0$ and simulate data with the following parameters:

T	sims	ϕ	θ	σ_η^2	σ_ε^2	β_0	β_1	β_2
1000	1	0	0.5	5	10	-100	0.01	3

Table 7: Parameter values for simulated REMARCS model with a latent MA(1) process





Figure 12: Simulated REMARCS Model Before and After Fitting

We then initiate the algorithm as usual and obtain the following parameter estimates:


Parameter	β_0	β_1	β_2	σ_ε^2	ϕ	θ	σ_η^2
True Value	-100.0000	0.0100	3.0000	10.0000	0.0000	0.5000	5.0000
Initial estimate	-100.0109	0.0103	2.9861	16.5249	0.0000	0.0000	0.0000
Final estimate	-100.0670	0.0104	2.9861	9.9510	-0.0816	0.5900	5.1878

Table 8: Parameter estimates for REMARCS model with latent MA(1) time series, before and after convergence

We take the final time series parameter estimates and perform inference on $\hat{\phi}$ and $\hat{\theta}_b$:



(a) Confidence region for (ϕ, θ_b)



(b) Confidence region for (ϕ, θ)

Figure 13: Confidence regions for observable and latent moving average parameters before restricting REMARCS model with $\theta_b = 0$




Figure 14: Confidence regions for observable and latent moving average parameters after restricting REMARCS model with $\theta_b = 0$

Note that while the point estimate $\hat{\theta}$ doesn't change much, the range of the confidence interval for θ has shrunk with the decrease in degrees of freedom.

We now look to the special case in which aggregate noise is an AR(1) process, which we recall can only occur if $\theta/\phi = \mathbb{E}(N^{-1})\sigma_\varepsilon^2/\sigma_\eta^2$. We select parameters for which this equality doesn't hold exactly, but is approximately true given the distribution of N_t :

T	sims	ϕ	θ	σ_η^2	σ_ε^2	β_0	β_1	β_2
1000	1	0.5	0.5	5.25	50	20	-0.001	2

Table 9: Parameter values for simulated REMARCS model with latent AR(1) process




Figure 15: Realization of a REMARCS Model in which the sum of the latent time series and aggregated cross-sectional noise is an AR(1) process

After fitting a REMARCS model to the data we obtain parameter estimates shown in Table 10 and fitted values in Figure 15.

Parameter	β_0	β_1	β_2	σ_ε^2	ϕ	θ	σ_η^2
True Value	20.0000	-0.0010	2.0000	50.0000	0.5000	0.5000	5.2500
Initial estimate	19.6845	-0.0007	2.0362	61.4985	0.0000	0.0000	0.0000
Final estimate	19.7719	-0.0008	1.9851	49.7772	0.4996	0.4962	5.1336

Table 10: Parameter estimates for REMARCS model with $\theta_b \approx 0$, before and after convergence

Note that while the point estimate for θ given by Table 10 is accurate, the confidence region thereof (Figure 16) practically spans all of $(0, 1)$, with which all that can be said is that θ is positive.




Figure 16: Inference on ARMA parameters of REMARCS model in which $\theta_b = 0$ with restricted model

Note in the left panel of the same figure we see that the estimate for θ_b is insignificant, prompting us to refit $\mathbf{w}$ with an AR(1) model and use H_{θ_b} (Equation ??) to construct a confidence region for $\hat{\theta}$, the products of which are shown in Figure 17:




Figure 17: Inference on ARMA parameters of REMARCS model in which $\theta_b = 0$ with restricted model

We now obtain a confidence interval for θ markedly smaller than before. Another change of note is that the correlation between $\hat{\phi}$ and $\hat{\theta}$ has gone from negative to positive. This makes sense given that the ratio between the two now has an added restriction.

Now that our algorithm has demonstrated promising results on simulated data, we observe how it performs on real RCS survey data.

5 Discussion

This dissertation develops a framework for analyzing high-frequency repeated cross-sectional survey data that exhibit temporal correlation while containing individual-level covariates that are time independent. It does so by introducing a latent ARMA process, inducing an overall ARMA-Noise residual structure, and recursively fitting parameter estimates with GLS and ARMA-Noise methods. In order to accommodate survey data with varying daily sample size, this research furthered ARMA-Noise modeling by relaxing the preexisting condition that sample sizes be constant. Results provide insight into how temporal correlation interacts with cross-sectional noise, particularly how not properly accounting for this interaction can obfuscate results and hinder proper inference. Given how large RCS datasets can be, computational methods were tailored to the specific characteristics of this model by leveraging the relationship between the spectral decompositions of the various covariance matrices involved. The exploration of these models has produced insight into the limitations of their analysis and yielded several interesting directions for further research.

A result central to this work is that parameter estimates for a latent ARMA process will be fundamentally flawed if fit to aggregate residuals, as the presence of additional white noise introduces bias in estimates for not only the variance parameter but the moving average parameter as well. Such an impact is tractable, however, and can be reverse engineered to obtain asymptotically unbiased estimates for the true time series parameters. The relation between the observable and latent time series must be incorporated into the analysis of time series estimate. This impact can change the order of the ARMA process on top of the parameters themselves. For instance, in the case that the latent time series is actually an AR(1) process, the observed aggregate errors will be ARMA(1,1) with a moving average parameter of opposite sign as that of the autoregressive parameter. Modeling the noisy time series as an AR(1) process would prevent a consistent estimate for ϕ from being obtained. Simulation results reinforce this conclusion. We see that as the relative variance of the additional noise increases, estimates for ϕ obtained by fitting an AR(1) model become increasingly biased, while those fit with an ARMA(1,1) model remain accurate.

This work provides a viable method for analyzing datasets such as the mVAM surveys, elucidating individual-level effects by properly accounting for temporal shocks that would otherwise obfuscate them. Such a method avoids the information loss inherent to the aggregate methods that are currently commonplace.

Furthermore, failure to account for temporal correlation when modeling repeated cross-sectional data can lead to invalid inference. At the same time, flattening the data to be analyzed into aggregates results in significant information loss at the individual level. In particular, in the event that individual-level

variables like education level or income source see little variance in their prevalence at the time-step level, fitting aggregated data using a time series with covariates could easily obscure an otherwise significant effect.

This method is not without limitations, particularly given its reliance on asymptotic distributions. As such confidence regions underperform for low values of T , particularly in boundary cases like $\phi, \theta \approx \pm 1$ and $\phi + \theta \approx 0$, as the resulting covariance structure is close to that of white noise. The imposition of an ARMA order no greater than $(1, 1)$ is certainly limiting, as it precludes REMARCS Models with all but three types of latent ARMA processes from being properly fit.

There are a number of natural extensions that arise from this work, particularly generalizing the time series level of the model to allow for $p, q > 1$. Given that survey data so often tracks binary output variables, relaxing the Gaussian condition is of further interest. Since confidence intervals are currently based on asymptotic distributions, the derivation of small-sample corrections is warranted.

References

- [1] George E. P. Box and Gwilym M. Jenkins. *Time Series Analysis: Forecasting and Control*. Holden-Day, San Francisco, CA, second edition, 1976.
- [2] Peter J. Brockwell and Richard A. Davis. *Time Series: Theory and Methods*. Springer Series in Statistics. Springer, New York, 2 edition, 1991.
- [3] Thomas A. DiPrete and David B. Grusky. Multilevel analysis of trends with an application to the study of income inequality in the united states. *American Sociological Review*, 55(3):438–451, 1990.
- [4] Wing keung Wong and Robert B. Miller. Repeated time series analysis of arima-noise models. *Journal of Business & Economic Statistics*, 8(2):243–250, 1990.
- [5] Lucia Modugno, Silvia Cagnone, and Simone Giannerini. A multilevel model with autoregressive components for the analysis of tribal art prices. *Journal of Applied Statistics*, 42(10):2141–2158, 2015.
- [6] A. J. Scott and T. M. F. Smith. Analysis of repeated surveys using time series methods. *Journal of the American Statistical Association*, 69(347):674–678, 1974.

6 Appendix

List of Tables

1	Implications of Various Restrictions of Time Series Parameters	8
---	--	---

2	Parameter values for simulated ARE model	12
3	Parameter values for simulated REMARCS model with all nonzero parameters	16
4	Parameter estimates for REMARCS model with all nonzero parameters, before and after convergence	18
5	Parameter values for simulated REMARCS model with a latent AR(1) process	18
6	Parameter estimates REMARCS model with latent AR(1) Time Series, before and after convergence	19
7	Parameter values for simulated REMARCS model with a latent MA(1) process	20
8	Parameter estimates for REMARCS model with latent MA(1) time series, before and after convergence	21
9	Parameter values for simulated REMARCS model with latent AR(1) process	22
10	Parameter estimates for REMARCS model with $\theta_b \approx 0$, before and after convergence	23

List of Figures

1	Flowchart of algorithm steps for estimating model parameters . .	5
2	Decision tree for selecting parameters to include in ARMA-Noise model	9
3	Realization of ARE model with parameters defined in Table 2 . .	12
4	Rates at which point estimates for τ_o fall within confidence regions based on true values	14
5	Rates at which point estimates for τ_l fall within confidence regions based on sample estimates	15
6	Empirical confidences with which model parameters fall within confidence regions generated after model convergence	16
7	A realization of a REMARCS model before and after fitting . . .	17
8	Confidence regions for time series parameter estimates of latent ARMA(1,1) process	17
9	Observations and fitted values from a REMARCS model with a latent AR(1) time series	18
10	Confidence regions for time series parameter estimates of latent ARMA(1,1) process	19
11	Confidence interval for ϕ after restricting model	20
12	Simulated REMARCS Model Before and After Fitting	20
13	Confidence regions for observable and latent moving average parameters before restricting REMARCS model with $\theta_b = 0$	21
14	Confidence regions for observable and latent moving average parameters after restricting REMARCS model with $\theta_b = 0$	22
15	Realization of a REMARCS Model in which the sum of the latent time series and aggregated cross-sectional noise is an AR(1) process	23

16	Inference on ARMA parameters of REMARCS model in which $\theta_b = 0$ with restricted model	24
17	Inference on ARMA parameters of REMARCS model in which $\theta_b = 0$ with restricted model	24