

Abstract

1 Introduction

Mobile phone surveys are an increasingly popular method for tracking household food security that is employed by aid organizations like the World Food Programme. Improving sampling capabilities and enabling real-time tracking of consumption and coping metrics, these surveys are conducted daily in an array of countries. Drawn with a random sampling strategy to select participants, samples do not necessarily overlap day to day, but are conducted at a frequency that often results in temporal correlation in the collected responses that hinders inference when using linear regression. While there already exist methods for analyzing data in which samples are surveyed repeatedly or aggregated at the time-step level, models that account for individual-level effects and population-level shocks are scarce. Those that do exist model temporal fluctuations with a latent time series that is recursively estimated and used to update parameter estimates via an EM algorithm, but these methods do not properly account for how cross-sectional noise impacts analysis. We correct this shortcoming by incorporating methods for analyzing data that are the result of multiple noisy observations of each element of a time series, known as an ARMA-Noise model.

We go on to establish that the use of an EM algorithm that factors the likelihood equation by conditioning on an estimate of a latent ARMA process incorrectly assumes that the resulting deviations will have a covariance structure conducive to standard fixed effect regression. We therefore replace the EM algorithm with one that recursively estimates variance parameters with an ARMA-Noise model and uses them to calculate GLS estimators for regression parameters. Confidence regions for time series parameters are generated via the Delta Method due to their latent nature. The respective implications of varying orders of the latent time ARMA process inform a decision rule for model selection.

While the newly incorporated method for fitting time series parameters corrects the bias that previously existed, it imposes a constant sample size at the time-step level, an unrealistic requirement in real-life surveying contexts. We therefore explored the extent to which ARMA-Noise modeling methods could be generalized to accommodate varying sample sizes.

With the proposed method fully developed, we test on simulated data and illustrate how model selection might be performed on a realization from a restricted model. We also include a case study on household survey data from Honduras that shows the method's potential to decorrelate residuals. The results of our analysis support the method as a viable option for modeling RCS survey data related to food security and other contexts.

2 Background

Repeated cross-sectional (RCS) data has long been the subject of analysis. When cross-sections are taken with high frequency, the possibility of temporal correlation arises. While panel data consist of a single sample being observed repeatedly, RCS data do not necessarily have any repeated observations, ruling out the treatment of the data as a multivariate time series. It is often instead the population mean that is treated as a time series, with individual observations modeled as a realization of that mean with some added noise.

Among the earliest methods designed for RCS data was that proposed by Scott and Smith [6], in which dependent variables are aggregated at the time step level and modeled as a time series with added noise:

$$y_t = \theta_t + e_t, \quad \text{Var}(e_t) = s_t^2 \quad (2.1)$$

In this model, surveys are conducted at regular time intervals and θ_t represents a population mean changing over time. The analytical method proposed therein introduces different possibilities for the covariance structures of $\{\theta_t\}$ and $\{e_t\}$, but is ultimately concerned with the use of all previously observed aggregates $\{y_t\}_{t=1}^T$ to estimate θ_T . As our objective is to improve the estimation and of covariate effects and their standard errors, we must look to more highly parameterized models.

Such a model was developed by DiPrete and Grusky [3], in which added complexity and flexibility is added to a standard fixed effects model by having the effects of individual-level variables depend on macro-level variables:

$$y_{ti} = X_{ti}\beta_t + \varepsilon_{ti}, \quad \beta_t = \mathbf{Z}_t\gamma, \quad \varepsilon_{ti} \sim \text{WN}(0, \sigma^2) \quad (2.2)$$

Such a model requires the presence of macro-level variables that vary over time and explain changing effects of individual-level variables. While this formulation has a broad range of applications, it is only appropriate in cases where temporal correlation can be assumed to be a function of known time-varying predictors, rather than an independent random process.

There also exist aggregate methods that involve calculating averages of independent and dependent variables at each time step and analyze these averages as a time series with covariates. Such a model is an example of a state-space model that can be modeled jointly via Kalman filtering [2]. In the regression setting, aggregates would be needed for both response and explanatory variables, fitting the following model:

$$\bar{\mathbf{y}} = (n^{-1}X'_tX_f)\beta + \mathbf{u}, \quad \mathbf{u} \sim \text{ARMA}(p, q)$$

While this method is effective at estimating population-level fixed effects like intercepts and slopes, it is limited when it comes to the estimation of individual-level effects, especially if indicator variables vary little in their distribution at each time step. This is a highly plausible scenario when it comes to sampling methods for surveys, as researchers interested in the effect of a certain covariate

may very well aim to be consistent in the proportions of time step samples that have this feature.

One model that has been developed to analyze RCS data with autocorrelated errors is the Autoregressive Random Effects (ARE) Model [5]. Treating observations as the sum of a fixed effects model and an AR(1) time series, ARE models have two levels: a fixed-effects level and a time series level.

The method employed to fit such a model is a recursive process that splits the log likelihood function of the model into two terms by conditioning on the latent time series $\mathbf{u}$. An EM algorithm is then employed for parameter estimation.

This conditional splitting is convenient, as $\mathbf{u} \sim N(\mathbf{0}, \mathbf{\Gamma})$ and $\mathbf{y}|\mathbf{u} \sim N(X_f\beta + \mathbf{u}, \sigma_\varepsilon^2 I_N)$. Furthermore, the choice of imposing an AR(1) structure on $\mathbf{u}$ lends itself to a concise representation of the determinant and inverse of $\mathbf{\Gamma}$, the covariance matrix of $\mathbf{u}$.

In the absence of an explicit observation of $\mathbf{u}$, its conditional expectation is calculated with current parameter estimates, then used to update those estimates. This process is repeated until convergence criteria are met.

There are two flaws in this methodology. The first is that fixed effect parameter estimates are updated with the assumption that $\mathbf{y} - \hat{\mathbf{u}}$ is a realization of a fixed effects model, when its covariance structure is altered by the error in estimating $\mathbf{u}$. The second problem stems from the effect of added individual noise on an AR(1) process, the sum of which forms an ARMA-Noise Process [4].

Let's assume the best case scenario for level one of the model, namely the perfect estimation of fixed effect parameters ($\hat{\beta} = \beta$). In this unlikely event, we are able to extract data consisting of multiple noisy observations of each element of an AR(1) process

With no prior information about the latent time series, the best estimates for the elements of $\mathbf{u}$ are the averages at each time step. The added noise turns an originally AR(1) process and turns it into an ARMA(1,1) process. Failing to add a moving average parameter necessarily introduces bias to the estimation of the time series parameters.

From Box (1976) [1], we have the following theorem:

If $\{y_t\}$ follows an ARMA(p_1, q_1) model, $\{a_t\}$ follows an ARMA(p_2, q_2) and the two series are independent, then their sum $w_t = y_t + a_t$ follows an ARMA(p, q) model in which $p \leq p_1 + p_2$ and $q \leq \max(p_1 + q_2, p_2 + q_1)$.

Because $\{a_t\}$ is white noise, $p_2 = q_2 = 0$. As $p_1 = 1$ and $q_1 = 0$, we can deduce that $\{w_t\}$ is an ARMA(p, q) process where $p \leq 1$ and $q \leq 1$. Because $\gamma_{\hat{\mathbf{u}}}(h+1)/\gamma_{\hat{\mathbf{u}}}(h) = \phi$ for $h > 0$, we know that $p = 1$. The ACVF of $\hat{\mathbf{u}}$ is therefore incompatible with that of an AR(1) process, so it must be ARMA(1,1).

As the best-case scenario of estimating and extracting the latent time series will still yield temporal correlation in the data, albeit reduced, we must turn to a more nuanced treatment of residual covariance. A natural choice is the GLS estimator of β , which we will see challenges the necessity of using $\hat{\mathbf{u}}$ in the estimation of β at all.

When framed in the context of repeated surveys, the requirement that the sample size at each time step is constant is unrealistic. Given that real life

survey data is rarely so uniform, generalizing the method to allow for varying daily sample size is a natural goal. In examining the method as it stands, the condition of fixed sample size is only necessary in the fitting of ARMA-Noise models, as $\mathbf{w} = \mathbf{u} + \mathbf{a}$ is only an ARMA process if $\{a_t = n_t^{-1} \sum_{i=1}^{n_t}\}$ is white noise, something that at first glance is violated by varying n_t . However, we will now show how this method can be adapted to allow for varying sample sizes, so long as certain conditions are met.

3 Method

Having identified what parts of existing methods we want to utilize as well as those that we wish to change or generalize, we are ready to define our model and accompanying algorithm, *REgression Models for Autocorrelated Repeated Cross-Sections*, or REMARCS:

3.1 Model and Algorithm

We define the vector of observations $\mathbf{y}$ as the realization of a model defined as the sum of fixed effects, a latent ARMA(1,1) process, and i.i.d. noise (*n.b. sample sizes are currently constant across time steps*).

$$y_{ti} = \mathbf{x}_{ti}\beta + u_t + \epsilon_{ti}, \quad i = 1, \dots, n \quad \epsilon_{ti} \sim \text{i.i.d. } N(0, \sigma_\epsilon^2) \quad (3.1)$$

$$u_t - \phi u_{t-1} = \eta_t + \theta \eta_{t-1}, \quad t = 1, \dots, T \quad \eta_t \sim N(0, \sigma_\eta^2) \quad (3.2)$$

The algorithm we propose is an iterative process applying the estimates of parameters from one level to update those of the other. The process is as follows:

- 1) Initialize with $\hat{\beta} = (X_f' X_f)^{-1} X_f' \mathbf{y}$
- 2) Estimate $(X_t \mathbf{u} + \varepsilon)$ with $(\mathbf{y} - X_f \hat{\beta})$ and fit ARMA-Noise model to obtain estimates $(\hat{\phi}, \hat{\theta}, \hat{\sigma}_\eta^2, \hat{\sigma}_\epsilon^2)$
- 3) Calculate $\hat{\Omega}$ and $l(\hat{\xi})$
- 4) Refit fixed effect parameter estimates by calculating GLS estimate $\hat{\beta}_{GLS} = (X_f' \hat{\Omega}^{-1} X_f) X_f' \hat{\Omega}^{-1} \mathbf{y}$
- 5) Repeat Steps 2-4 until convergence criteria are met
- 6) Generate confidence regions. If any parameters are insignificant, restrict model and repeat algorithm

3.2 Evaluating Convergence

Instead of splitting the likelihood by conditioning on $\hat{\mathbf{u}}$, doing so with $\bar{\mathbf{y}}$ instead of has the advantage that likelihoods are based on a constant statistic.

Let's recall the log likelihood $l(\xi)$:

$$l(\xi) = -\frac{N}{2}\ln(2\pi) - \frac{1}{2}\ln|\Omega| - \frac{1}{2}(\mathbf{y} - X_f\beta)' \Omega^{-1}(\mathbf{y} - X_f\beta)$$

We now condition on $\bar{\mathbf{y}} := n^{-1}X'_t\mathbf{y}$, which has the following distribution.

$$\bar{\mathbf{y}} \sim N(\mu_{\bar{\mathbf{y}}}, \Sigma_{\bar{\mathbf{y}}}), \quad \mu_{\bar{\mathbf{y}}} = n^{-1}X'_tX_f\beta, \quad \Sigma_{\bar{\mathbf{y}}} = \Gamma + n^{-1}\sigma_\varepsilon^2 I_T$$

Defining ξ to be the vector of parameters $(\beta, \phi, \theta, \sigma_\eta^2, \sigma_\varepsilon^2)$, we can then define the log likelihood function:

$$l_t(\xi) = \ln[f(\bar{\mathbf{y}}|\xi)] = -\frac{T}{2}\ln(2\pi) - \frac{1}{2}\ln|\Sigma_{\bar{\mathbf{y}}}| - \frac{1}{2}(\bar{\mathbf{y}} - \mu_{\bar{\mathbf{y}}})' \Sigma_{\bar{\mathbf{y}}}^{-1}(\bar{\mathbf{y}} - \mu_{\bar{\mathbf{y}}})$$

If we then condition $\mathbf{y}$ on $\bar{\mathbf{y}}$, the only randomness left is in the vector of deviations, which are independent of $\bar{\mathbf{y}}$:

$$\mathbf{y} - X_t\bar{\mathbf{y}} = \varepsilon - X_t\mathbf{a}$$

This is a vector with mean $\mathbf{0}$ and variance:

$$\text{Var}(\varepsilon - X_t\mathbf{a}) = \sigma_\varepsilon^2(I_N - n^{-1}X_tX'_t) \quad (3.3)$$

We now derive the second part of our split likelihood. Because of the rank-deficient nature of the matrix Equation 3.3, we can simplify the log likelihood

$$l_i(\xi) = \ln f(\mathbf{y} | \bar{\mathbf{y}}, \sigma_\varepsilon^2) = -\frac{N-T}{2}\ln(2\pi\sigma_\varepsilon^2) - \frac{1}{2\sigma_\varepsilon^2} \sum_{t=1}^T \sum_{i=1}^n (y_{it} - \bar{y}_t)^2, \quad (3.4)$$

With this we can express the full likelihood as $l(\xi) = l_t(\xi) + l_i(\sigma_\varepsilon^2)$. However, since $\bar{\mathbf{y}}_t$ is independent of our model selection and parameter estimates, l_i is going to be relatively stable. We therefore pair $\xi \in \mathbb{R}^k$ and $l_t(\xi)$ to calculate the Akaike Information Criterion (AIC):

$$AIC = -2k - 2l_t(\xi)$$

Calculating this weighted likelihood is how we will compare the performance of models that include different subsets of parameters.

3.3 Inference

Once parameter estimates have converged, confidence intervals are constructed.

Fixed effect estimates have covariance matrix:

$$\text{Var}(\hat{\beta}_{GLS}) = (X_f' \Sigma_{\mathbf{y}}^{-1} X_f)^{-1} \quad (3.5)$$

This covariance matrix is estimated with final parameter estimates and then used to construct confidence regions. Just as this method of building confidence regions for β in some sense assumes that parameters $\phi, \theta, \sigma_\eta^2, \sigma_\varepsilon^2$ have been perfectly estimated, confidence intervals for those same parameters makes an analogous assumption.

If β is estimated perfectly, then $\mathbf{y} - X_f \hat{\beta} = X_t \mathbf{u} + \varepsilon =: \mathbf{z}$, and time-step averages $w_t = \sum_{i=1}^n z_{ti}$ are equal to the sum $u_t + a_t$, with $a_t \stackrel{i.i.d.}{\sim} N(0, \sigma_\varepsilon^2/n)$. According to Box and Jenkins [1], if $\mathbf{u}$ is an ARMA(p_u, q_u) process with $p_u, q_u \leq 1$, then $\mathbf{w}$ is an ARMA(p_w, q_w) process, with $p_w, q_w \leq 1$.

What's most important to note is that if $\mathbf{u}$ is parameterized by $\phi, \theta, \sigma_\eta^2$, then $\mathbf{w}$ is parameterized by $\phi, \theta_b, \sigma_b^2$. According to [2], the Maximum Likelihood estimates for $\phi, \theta_b, \sigma_b^2$, which can be obtained using an R package like *arima*, have the following asymptotic distribution (when normalized):

$$\sqrt{T} \left[\begin{pmatrix} \hat{\phi} \\ \hat{\theta}_b \\ \hat{\sigma}_b^2 \end{pmatrix} - \begin{pmatrix} \phi \\ \theta_b \\ \sigma_b^2 \end{pmatrix} \right] \xrightarrow{d} N(\mathbf{0}, \Sigma_b)$$

$$\Sigma_b = \begin{bmatrix} \frac{(1-\phi^2)(1+\phi\theta_b)^2}{(\phi+\theta_b)^2} & \frac{-(1-\theta_b^2)(1-\phi^2)(1+\phi\theta_b)}{(\phi+\theta_b)^2} & 0 \\ \frac{-(1-\theta_b^2)(1-\phi^2)(1+\phi\theta_b)}{(\phi+\theta_b)^2} & \frac{(1-\theta_b^2)(1+\phi\theta_b)^2}{(\phi+\theta_b)^2} & 0 \\ 0 & 0 & 2\sigma_b^4 \end{bmatrix}$$

The parameter estimate $\hat{\sigma}_\varepsilon^2$ and its normalized asymptotic variance are given by:

$$\hat{\sigma}_\varepsilon^2 = (N - T)^{-1} \sum_{t=1}^T \sum_{i=1}^{n_t} (w_{ti} - \bar{w}_t.)^2 \quad (3.6)$$

$$\sqrt{T}(\hat{\sigma}_\varepsilon^2 - \sigma_\varepsilon^2) \xrightarrow{d} N(0, \frac{2\sigma_\varepsilon^4}{n-1})$$

Because $y_t. \perp S_t^2$ for all t , we know that $\hat{\sigma}_\varepsilon^2 \perp (\hat{\phi}, \hat{\theta}_b, \hat{\sigma}_\eta^2)$ and can give the following asymptotic distribution for “observable” parameters $\tau_o = (\phi, \theta_b, \sigma_b^2, \sigma_\varepsilon^2)$:

$$\sqrt{T}(\hat{\tau}_o - \tau_o) \xrightarrow{d} N(\mathbf{0}, \Sigma_o), \quad \Sigma_o = \begin{bmatrix} \Sigma_b & \mathbf{0} \\ \mathbf{0} & \frac{2\sigma_\varepsilon^2}{n-1} \end{bmatrix} \quad (3.7)$$

Equipped with the above estimates, Wong and Miller detail how to estimate “latent” parameters $\tau_l = (\theta, \sigma_\eta^2)$ by relating the autocovariance functions of $\mathbf{u}$ and $\mathbf{w}$ through the system of equations $H(\theta, \sigma_\eta^2)$, which is obtained by taking

the expression $w_t = u_t + a_t$, multiplying $(1 - \phi B)$, with B symbolizing the backwards operator, then calculating the ACVF of both sides at lags 0 and 1. This system of equations is then supplemented by $\phi - \phi_*$ and $\sigma_\varepsilon^2 - \sigma_{\varepsilon*}^2$ so the final covariance matrix includes all parameters of interest.

$$H(\theta, \sigma_\eta^2) = \begin{cases} \phi - \phi_* \\ (1 + \theta_b^2)\sigma_b^2 - (1 + \theta^2)\sigma_\eta^2 - (1 + \phi^2)\sigma_a^2 \\ \theta_b\sigma_b^2 - \theta\sigma_\eta^2 + \phi\sigma_a^2 \\ \sigma_\varepsilon^2 - \sigma_{\varepsilon*}^2 \end{cases} \quad (3.8)$$

$$\tau_0 = \begin{pmatrix} \phi \\ \theta_b \\ \sigma_b^2 \\ \sigma_\varepsilon^2 \end{pmatrix} \quad \tau_1 = \begin{pmatrix} \phi_* \\ \theta \\ \sigma_\eta^2 \\ \sigma_{\varepsilon*}^2 \end{pmatrix} \quad (3.9)$$

Estimates for τ_l are then given by the roots of H . It is important to note that estimates for τ_o can produce a system of equations that has no real roots, especially for relatively low values of T . In this event we use numerical methods to solve for parameter estimates $\hat{\theta}, \hat{\sigma}_\eta^2$ that minimize $\|H\|^2$.

Were we to have an explicit function $(\hat{\phi}, \hat{\theta}, \hat{\sigma}_\eta^2) = F(\hat{\tau}_o)$ providing estimates for τ_l , we would directly apply the Delta Method to obtain the asymptotic distribution of our vector of parameter estimates:

$$\sqrt{T}(F(\hat{\tau}_o) - \tau_l) \xrightarrow{d} N(\mathbf{0}, G\Sigma_o G'), \quad G = \frac{\partial F}{\partial(\tau_o)} \quad (3.10)$$

Because we're using H , a function whose roots are our estimates, we will adapt 3.10 using the Implicit Function Theorem:

$$G = - \left(\frac{\partial H_f}{\partial \tau_1} \right)^{-1} \left(\frac{\partial H}{\partial \tau_0} \right) \quad (3.11)$$

Because we're ultimately interested in the normalized asymptotic distribution of parameter estimates $(\hat{\phi}, \hat{\theta}, \hat{\sigma}_\eta^2, \hat{\sigma}_\varepsilon^2)$, we expand H to include some identity functions.

Our procedure is to fit $\hat{\tau}_o$, solve for $\hat{\tau}_l$, and then plug in all estimates into Equations 3.10 and 3.11. If any parameter estimates, observed or latent, are deemed insignificant, the relevant parameter can be restricted to 0 and the procedure adapted accordingly. This decision procedure is detailed next.

3.4 Model Selection

Before outlining the methodology for model selection, we motivate its design by exploring the implications of certain parameter combinations in an ARMA-Noise model $\mathbf{w} = \mathbf{u} + \mathbf{a}$.

From these relations we derive a decision rule for model selection:

Parameter Restriction	Order of $\mathbf{u}$	Order of $\mathbf{w}$
$\theta = 0$	AR(1)	ARMA(1,1)
$\phi = 0$	MA(1)	MA(1)
$\frac{\theta}{\phi} = \frac{\sigma_\varepsilon^2}{n\sigma_\eta^2}$	ARMA(1,1)	AR(1)
$\phi = \theta = 0$	WN	WN
Else	ARMA(1,1)	ARMA(1,1)

Table 1: Implications of Various Restrictions of Time Series Parameters




Figure 1: Decision tree for selecting parameters to include in ARMA-Noise model

3.5 Generalizing to Varying Sample Size

Having detailed the steps involved in fitting and performing inference on model parameter estimates, we now turn to the condition required for fitting an ARMA-Noise model, namely that sample size is constant at each time step ($n_t = n \forall t$). This condition is included so that the time series $\{a_t\}$ is Gaussian white noise, and $\mathbf{w}$ is therefore still an ARMA process.

Given that RCS data is rarely of a fixed sample size, we now explore the extent to which this condition can be relaxed whilst $\mathbf{a}$ remains white noise. We will now treat $\{n_t\}$ as the realizations of random variables $\{N_t\}_{t=1}^T$. What we aim to establish are the conditions under which the random variable $A_t = N_t^{-1} \sum_{i=1}^{N_t} \varepsilon_{ti}$ is still white noise.

Claim:

Let $\{\varepsilon_{ti}\}$ be i.i.d. random variables with

$$\mathbb{E}(\varepsilon_{ti}) = 0, \quad \text{Var}(\varepsilon_{ti}) = \sigma_\varepsilon^2 < \infty,$$

Let $\{N_t\}$ be a stochastic process such that

$$P(N_t \in \mathbb{N} \setminus \{0\}) = 1 \quad \mathbb{E}(N_t^{-1}) \quad \text{for all } t,$$

Define

$$A_t := \frac{1}{N_t} \sum_{i=1}^{N_t} \varepsilon_{ti}.$$

Then $\{A_t\}$ is white noise with $\sigma_a^2 = \sigma_\varepsilon^2 \mathbb{E}(N_t^{-1})$

Proof:

First, let $\{N_t\}$ be i.i.d. random variables with a sample space that is a subset of the integers strictly greater than 0. The conditional random variable $A_t|N_t$ then has variance σ_ε^2/N_t .

To find the marginal variance $\text{Var}(A_t)$, we use the Law of Total Variation:

$$\text{Var}(A_t) = \mathbb{E}(\text{Var}(A_t|N_t)) + \text{Var}(\mathbb{E}(A_t|N_t)) \quad (3.12)$$

Now, $\text{Var}(A_t|N_t) = \sigma_\varepsilon^2 N_t^{-1}$, and therefore has an expectation of $\sigma_\varepsilon^2 \mathbb{E}(N^{-1})$, which is well defined given that the sample space of N_t is a subset of the integers strictly greater than zero. At the same time, the expectation of A is 0 regardless of the value of N , meaning the second term in Equation 3.12 is zero, leaving only $\text{Var}(A) = \sigma_\varepsilon^2 \mathbb{E}(N^{-1})$.

Proof:

$$N \geq 1 \quad \text{a.s.}$$

$$\Rightarrow 0 < N^{-1} \leq 1 \quad \text{a.s.}$$

$$\Rightarrow 0 < \mathbb{E}(N^{-1}) \leq 1 \quad \text{a.s.}$$

We can quickly verify that A_t is zero mean with the Law of Total Expectation:

$$\mathbb{E}(A_t) = \mathbb{E}(\mathbb{E}(A_t|N_t)) = \mathbb{E}(0) = 0$$

Now we've already established $\text{Var}(A_t)$ to be independent of t , so we now verify that $\text{Cov}(A_r, A_s) = 0$ for $r \neq s$. The Law of Total Covariance states:

$$\text{Cov}(A_r, A_s) = \mathbb{E}(\text{Cov}(A_r, A_s|N_r, N_s)) + \text{Cov}(\mathbb{E}(A_r|N_r), \mathbb{E}(A_s|N_s))$$

Given that A_r and A_s are functions of independent random vectors $\varepsilon_r \in \mathbb{R}^{n_r}$ and $\varepsilon_s \in \mathbb{R}^{n_s}$, $\text{Cov}(A_r, A_s|N_r, N_s)$ must be zero. Having already established that both variables are zero-mean, the above equation reduces to:

$$\text{Cov}(A_r, A_s) = \mathbb{E}(0) + \text{Cov}(0, 0) = 0$$

□

Having shown that $\{A_t\}$ is still white noise, it now remains to estimate its variance. Though the sample sizes now vary, our estimate for σ_ε^2 hasn't changed (see Equation 3.6).

Without requiring the distribution of N_t , σ_a^2 can be estimated with:

$$\hat{\sigma}_a^2 = \frac{\hat{\sigma}_\varepsilon^2}{T} \sum_{t=1}^T \frac{1}{N_t} \quad (3.13)$$

With $\hat{\sigma}_\varepsilon^2$, $\hat{\sigma}_a^2$ and $\hat{\text{Var}}(\hat{\sigma}_\varepsilon^2)$ defined as above, the algorithm proposed by Wong and Miller can be once again applied, with necessary adjustments made to Σ_o and H .

The takeaway from all this is that a constant expected inverse of $\{N_t\}$ is the critical piece for $\{A_t\}$ to still be white noise.

So long as the set of inverse sample sizes $\{n_t^{-1}\}$ is conceivably of constant expectation, the ARMA-Noise method is applicable.

COMMENT ON ASSUMPTION OF INDEPENDENCE IN CALCULATING VARIANCE OF ESTIMATE AND IMPLICATIONS FOR STRICTLY STATIONARY SAMPLE SIZES AND NON-CONSTANT MEAN

MATRIX N and how the likelihood changes

Comment on missing time steps

3.6 Extending to Mixed Effects Models

A natural extension of this algorithm is to allow for random effects in addition to fixed ones, yielding a model of the following form:

$$\mathbf{y} = X_f \beta + X_r \mathbf{b} + X_t \mathbf{u} + \varepsilon, \quad \text{Var}(\mathbf{b}) = \Sigma_b \quad (3.14)$$

The algorithm changes only in that the initial model fit includes random effects (*lmer package*) The fixed effects estimate $\hat{\beta}_{GLS}$ adjusts in that Ω now

includes the term $X_r \Sigma_b X_r'$. Random intercepts $\mathbf{b}$ are updated with the conditional expectation $\mathbb{E}(\mathbf{b}|\mathbf{y}, \hat{\xi}) = \Sigma_b X_r' \Omega^{-1}(\mathbf{y} - X_f \hat{\beta})$. Those intercepts are then used to update the parameters of Σ_b , using the formula INSERT FORMULA HERE. A flowchart depicting the steps of the algorithm is given in Figure 2.




Figure 2: Flowchart of algorithm steps for estimating model parameters

MAKE SURE ALL NOTATION HAS BEEN DEFINED!!!!

4 Simulations

Before applying this methodology to real data, we first corroborate its necessity and efficacy through simulations, closely examining individual components of the algorithm along with pre-existing methods.

4.1 Effect of Noise on ARMA Processes

We begin by exploring how observation noise impacts parameter estimation for AR(1) processes. Take the following model, in which white noise is added to such a process:

$$w_t = u_t + a_t, \quad u_t - \phi u_{t-1} = \eta_t \quad (4.1)$$

$$a_t \stackrel{i.i.d.}{\sim} N(0, \sigma_a^2), \quad \eta_t \stackrel{i.i.d.}{\sim} N(0, \sigma_\eta^2)$$

For varying values of ϕ and fixed parameters ($\sigma_a^2 = 2, \sigma_\eta^2 = 1, T = 1000$), 130 realizations were simulated, each being fit with an AR(1) model and an ARMA(1,1) model. Figure 3 displays the aggregate means for both models. It is immediately apparent that ϕ is only estimated effectively by the latter model, supporting the claim that, while $\mathbf{u}$ may be AR(1), $\mathbf{w}$ is an ARMA(1,1) process.




Figure 3: Aggregated estimates for AR. An interactive version of this figure is available [here](#).

It is also worth noting that it is not currently possible to tease apart the two sources of noise. This is only possible if there are multiple noisy observations of each time series element. We now proceed to examine the efficacy with which the methodology of fitting ARMA-Noise Models—here generalized for varying sample sizes—does just that.

4.2 Empirical Confidence of ARMA-Noise Models

Before applying the REMARCS algorithm in its entirety to simulated data, it would be prudent to verify that generalizing ARMA-Noise Models to allow for varying sample sizes still produces valid estimates.

We begin with simulations of an ARMA(1,1)-Noise Model:

$$y_{ti} = u_t + \varepsilon_{ti}, \quad u_t - \phi u_{t-1} = \eta_t + \theta \eta_{t-1}$$

$$t = 1, \dots, T \quad i = 1, \dots, n_t \quad n_t \stackrel{i.i.d.}{\sim} \text{Poi}(\lambda) + 1$$

Differing values of ϕ and θ were combined with fixed parameters $\sigma_\varepsilon^2 = 1$, $\sigma_\eta^2 = 1$, and $\lambda = 5$. For each value-pair (ϕ, θ) , 10 simulations are generated and used to estimate parameter values. Covariance matrices are estimated and used to calculate empirical confidence levels with $\alpha = 0.05$ (see Figure 4).

COMMENT on use of method of moments estimator for wong

defined by $w_t = \bar{\mathbf{y}}_t$, with MA parameter θ_b and innovation variance σ_b^2 , produces time series estimates with normalized covariance close to Σ_{ϕ, θ_b} . For each set of parameters, we calculate the true values of θ_b and σ_b^2 . We then fit $(\hat{\phi}, \hat{\theta}_b)$ with $\bar{\mathbf{w}}$ and construct confidence regions centered around those estimates, counting the rates at which these confidence regions capture the true parameter values. These empirical confidences are displayed in Figure 4.

CHECK THAT DELTA METHOD IS TAKING DERIVATIVES W.R.T. VEPS NOT VARA




Figure 4: Rates at which point estimates for τ_l fall within confidence regions based on sample estimates

We can see performance dipped for parameter combinations (ϕ, θ) close to the $\phi + \theta = 0$ line for time series of length $T = 1000$. However, we can see that performance improved for $T = 5000$, suggesting that parameter estimates for ARMA(1,1) processes that resemble white noise take longer to converge to their asymptotic distributions.

4.3 Instability of EM Algorithm

To highlight the necessity for a different approach to fitting latent parameters, we test the stability of time series parameters using an EM algorithm. We define our model to be as follows:

$$y_{ti} = \mathbf{x}_{ti}\beta + u_t + \varepsilon_{ti}, \quad t = 1, \dots, T$$

$$i = 1, \dots, n_t \quad n_t \stackrel{i.i.d.}{\sim} \text{Poi}(\lambda) + 1$$

Here, $\mathbf{u}$ is defined to be an AR(1) process with innovation variance $\sigma_\eta^2 = 1$. Other fixed parameters include $\sigma_\varepsilon^2 = 20$, $T = 1000$, and $\beta = c(43, -0.001, 0.25)$. For varying values of ϕ , 9 realizations of the model were simulated, and for each simulation, *true values for all parameters* were used to estimate $\hat{\mathbf{u}} = \Sigma_{\mathbf{u}} X_t' \Sigma_{\mathbf{y}}^{-1} (\mathbf{y} - X_f \beta)$. That estimate was then used to fit an AR(1) model, updating $\hat{\phi}$ just as would be done via the EM algorithm (in the unlikely event all parameters were unwittingly estimated perfectly). Results are displayed in Table 2.

Table 2: Aggregate Estimates and Empirical Confidence Levels for AR parameter Following Ideal Time Series Estimation

phi	0.10	0.30	0.50	0.70	0.90
phi_mean	0.18	0.48	0.73	0.89	0.97
phi_captured	0.22	0.00	0.00	0.00	

It is immediately apparent that even in the event of perfect parameter estimation, the resulting time series estimate $\hat{\mathbf{u}}$ is not an AR(1) process and more importantly will yield biased estimates for time series parameters if treated as one. Having demonstrated the need for the REMARCS Algorithm, we now turn to evaluating its efficacy.

4.4 Empirical Confidence of REMARCS algorithm

Now

In order to assess the efficacy of this method, 2 simulations were conducted with fixed parameters $\beta = (43, -0.001, 0.25)$, $\sigma_\eta^2 = 1$, $\sigma_\varepsilon^2 = 20$, and varying combinations of ARMA parameters (ϕ, θ) .

In Figure 5, we can see the proportion of simulations that yielded confidence regions successfully capturing β (left) and $(\phi, \theta, \sigma_\eta^2)$ (right). One observation we can make is that, similarly to the ARMA-Noise simulations, convergence to asymptotic distributions is slower for parameter combinations close to the line $\phi + \theta = 0$. This under-performance is mitigated as T grows, indicating that time series more closely resembling white noise will require larger values of T for accurate estimation.

EXCLUDE VETA FROM CONFIDENCE REGIONS FOR RIGHT PANEL




Figure 5: Empirical confidences with which model parameters fall within confidence regions generated after model convergence

4.5 Empirical Confidence of Model Selection

Now that we've seen some aggregate results displaying the efficacy of this method under different parameters, we aim to evaluate the algorithm's ability to identify when model restriction is necessary.

We start with the case in which $\mathbf{u}$ is an AR(1) process. We simulate REMARCS models with such a latent time series, tracking the rate at which the algorithm yields an insignificant estimate for θ . Fixed parameters include $T = 1000$, $\sigma_\eta^2 = 1$, $\sigma_\varepsilon^2 = 20$, $\lambda = 5$, $\beta = (43, -0.001, 0.25)$. For each value of ϕ , 2 realizations were simulated and the REMARCS algorithm applied. Displayed below are the rates at which parameter estimates for θ were deemed insignificant.

Table 3: Rates at which MA Parameter is Insignificant When True Parameter is Zero

phi	-0.90	-0.50	-0.10	0.30	0.70
theta_captured	1.00	0.50	1.00	0.50	1.00

We now proceed to the case in which $\mathbf{u}$ is an MA(1) process. We simulate REMARCS models with such a latent time series, tracking the rate at which the algorithm yields an insignificant estimate for ϕ . Fixed parameters include $T = 1000$, $\sigma_\eta^2 = 1$, $\sigma_\varepsilon^2 = 20$, $\lambda = 5$, $\beta = (43, -0.001, 0.25)$. For each value of θ , 2 realizations were simulated and the REMARCS algorithm applied. Displayed below are the rates at which parameter estimates for ϕ were deemed insignificant.

Table 4: Rates at which AR Parameter is Insignificant When True Parameter is Zero

theta	-0.90	-0.50	-0.10	0.30	0.70
phi_captured	1.00	1.00	0.50	0.50	1.00

From these results we can see that the algorithm effectively identifies when model restriction is necessary.

4.6 Method Comparison

We begin with a realization of a REMARCS model (Figure 6) with parameters detailed in Table 5. The moving average parameter θ has been set to zero so the model is consistent with the ARE model detailed by Modugno et al.

Our best attempt to recreate the ARE algorithm encounters convergence problems, likely due to inappropriately fitting time series estimates with an AR(1) time series.

Parameter	Intercept	Slope	Dummy Indicator	σ_ε^2	ϕ	θ	σ_η^2
Value	43	-0.001	0.25	20	0.9	0	1

Table 5: Parameter values for simulated ARE model




Figure 6: Realization of ARE model with parameters defined in Table 5

We now compare the performance of the REMARCS algorithm with those of the ARE model and fitting a time series with covariates.

We consider fitting this model by aggregating the data at each time point and treating the averages as a time series with covariates. 9 simulations were generated and fit with each method.

Some sample rows of simulated data are provided below:

t	x	y
1.00	1	45.76
1.00	1	46.86
1.00	0	36.00
1.00	1	39.28
1.00	0	42.10
1.00	1	39.93

After aggregating at each time step, we obtain the following condensed dataset, from which a time series with covariates is then fit:

t	n	x	y
1.00	7	0.71	42.01
2.00	6	0.50	39.34
3.00	8	0.62	39.61
4.00	6	0.50	43.04
5.00	7	0.43	38.24
6.00	5	0.40	43.79

If we then fit a time series with covariates to the condensed dataset, we obtain results with empirical means displayed in the table below. The “Significance Frequency” column displays the fraction of the hundred simulations in which the corresponding parameter estimate was statistically significant at the $\alpha = .05$ level.

Parameter	True Value	TSCOV Mean	REMARCS Mean	Sig. Freq. TS	Sig. Freq. REMARCS
phi	0.900	0.888	0.888	1.000	1.000
theta	0.000	-0.591	0.035	1.000	0.000
Intercept	43.000	42.932	42.915	1.000	1.000
Slope	-0.001	-0.001	-0.001	0.111	0.111
Covariate	0.250	0.273	0.264	0.222	0.667

Note that entries of the x column, while displaying some variability, are fairly close to 0.5.

ADD STANDARD ERRORS!!! CHANGE SIGNIFICANCE TO EMPIRICAL CONFIDENCE DOUBLE CHECK SIGMAO FUNCTION, VAREPS OR VARA

- SIMULATE FIXED/MIXED EFFECT REMARCS DATA
- FIT WITH FE/ME MODEL
- FIT WITH TS WITH COVARIATES
- MODUGNO INSTABILITY
- FIT WITH REMARCS

- HEATMAPS
- MODEL SELECTION

RCSFIXED CANT CURRENTLY HANDLE MISSING VALUES

Now that our algorithm has demonstrated promising results on simulated data, we observe how it performs on real RCS survey data.

5 Discussion

This dissertation develops a framework for analyzing high-frequency repeated cross-sectional survey data that exhibit temporal correlation while containing individual-level covariates that are time independent. It does so by introducing a latent ARMA process, inducing an overall ARMA-Noise residual structure, and recursively fitting parameter estimates with GLS and ARMA-Noise methods. In order to accommodate survey data with varying daily sample size, this research furthered ARMA-Noise modeling by relaxing the preexisting condition that sample sizes be constant. Results provide insight into how temporal correlation interacts with cross-sectional noise, particularly how not properly accounting for this interaction can obfuscate results and hinder proper inference. Given how large RCS datasets can be, computational methods were tailored to the specific characteristics of this model by leveraging the relationship between the spectral decompositions of the various covariance matrices involved. The exploration of these models has produced insight into the limitations of their analysis and yielded several interesting directions for further research.

A result central to this work is that parameter estimates for a latent ARMA process will be fundamentally flawed if fit to aggregate residuals, as the presence of additional white noise introduces bias in estimates for not only the variance parameter but the moving average parameter as well. Such an impact is tractable, however, and can be reverse engineered to obtain asymptotically unbiased estimates for the true time series parameters. The relation between the observable and latent time series must be incorporated into the analysis of time series estimate. This impact can change the order of the ARMA process on top of the parameters themselves. For instance, in the case that the latent time series is actually an AR(1) process, the observed aggregate errors will be ARMA(1,1) with a moving average parameter of opposite sign as that of the autoregressive parameter. Modeling the noisy time series as an AR(1) process would prevent a consistent estimate for ϕ from being obtained. Simulation results reinforce this conclusion. We see that as the relative variance of the additional noise increases, estimates for ϕ obtained by fitting an AR(1) model become increasingly biased, while those fit with an ARMA(1,1) model remain accurate.

This work provides a viable method for analyzing datasets such as the mVAM surveys, elucidating individual-level effects by properly accounting for temporal shocks that would otherwise obfuscate them. Such a method avoids the information loss inherent to the aggregate methods that are currently commonplace.

Furthermore, failure to account for temporal correlation when modeling repeated cross-sectional data can lead to invalid inference. At the same time, flattening the data to be analyzed into aggregates results in significant information loss at the individual level. In particular, in the event that individual-level variables like education level or income source see little variance in their prevalence at the time-step level, fitting aggregated data using a time series with covariates could easily obscure an otherwise significant effect.

This method is not without limitations, particularly given its reliance on asymptotic distributions. As such confidence regions underperform for low values of T , particularly in boundary cases like $\phi, \theta \approx \pm 1$ and $\phi + \theta \approx 0$, as the resulting covariance structure is close to that of white noise. The imposition of an ARMA order no greater than $(1, 1)$ is certainly limiting, as it precludes REMARCS Models with all but three types of latent ARMA processes from being properly fit.

There are a number of natural extensions that arise from this work, particularly generalizing the time series level of the model to allow for $p, q > 1$. Given that survey data so often tracks binary output variables, relaxing the Gaussian condition is of further interest. Since confidence intervals are currently based on asymptotic distributions, the derivation of small-sample corrections is warranted. Lastly, while the theory now accommodates samples size that vary, whether it be strictly stationary or not in certain circumstances, understanding the effect this has on parameter estimation is not yet understood.

References

- [1] George E. P. Box and Gwilym M. Jenkins. *Time Series Analysis: Forecasting and Control*. Holden-Day, San Francisco, CA, second edition, 1976.
- [2] Peter J. Brockwell and Richard A. Davis. *Time Series: Theory and Methods*. Springer Series in Statistics. Springer, New York, 2 edition, 1991.
- [3] Thomas A. DiPrete and David B. Grusky. Multilevel analysis of trends with an application to the study of income inequality in the united states. *American Sociological Review*, 55(3):438–451, 1990.
- [4] Wing keung Wong and Robert B. Miller. Repeated time series analysis of arima-noise models. *Journal of Business & Economic Statistics*, 8(2):243–250, 1990.
- [5] Lucia Modugno, Silvia Cagnone, and Simone Giannerini. A multilevel model with autoregressive components for the analysis of tribal art prices. *Journal of Applied Statistics*, 42(10):2141–2158, 2015.
- [6] A. J. Scott and T. M. F. Smith. Analysis of repeated surveys using time series methods. *Journal of the American Statistical Association*, 69(347):674–678, 1974.

6 Appendix

List of Tables

1	Implications of Various Restrictions of Time Series Parameters . . .	8
2	Aggregate Estimates and Empirical Confidence Levels for AR parameter Following Ideal Time Series Estimation	14
3	Rates at which MA Parameter is Insignificant When True Parameter is Zero	15
4	Rates at which AR Parameter is Insignificant When True Parameter is Zero	16
5	Parameter values for simulated ARE model	16

List of Figures

1	Decision tree for selecting parameters to include in ARMA-Noise model	8
2	Flowchart of algorithm steps for estimating model parameters . .	11
3	Aggregated estimates for AR. An interactive version of this figure is available here.	12
4	Rates at which point estimates for τ_l fall within confidence regions based on sample estimates	13
5	Empirical confidences with which model parameters fall within confidence regions generated after model convergence	15
6	Realization of ARE model with parameters defined in Table 5 . .	16